

Design Myths

Dejan Sarka

Introduction

- Dejan Sarka (dsarka@solidq.com, dsarka@siol.net, @DejanSarka)
 - 30 years of experience
 - SQL Server MVP, MCT,...
 - 15 books
 - 15+ courses
 - Focus:
 - Data modeling
 - Data mining
 - Data quality

Storage Myth

- Relational database is just a storage system
- WRONG!
- Relational database is the heart of IT systems
 - Relational model is about data integrity
 - A RDBMS enforces data integrity

3

Physical Myth

- Knowing details of a RDBMS is very important
 - Design for performance, not for theory
- WRONG!
- Logical problems are more important than physical ones
 - Physical problems pass over
 - Relational model is a logical model
 - A good logical design should perform well; the opposite is not always true

4

Repeating Group Myth

- A table with repeating group of columns needs normalization (1st NF)
 - EMPLOYEES(EmpId, EmpName, Child1, Child2, Child3,...)
- WRONG!
- Repeating columns with similar names is not a repeating group!
 - Just a constraint built in the model
 - ORDERS(OrdId, CustName, OrderDate, DueDate, ShipDate,...)

5

Atomicity Myth

- Every collection in a single column of a single row is a repeating group, is not atomic
- WRONG!
- We need to know the correct definition of atomicity
 - A value is atomic as long as you operate with it only with methods defined on the data type of the value (e.g. Datetime data)
 - XML and UDTs with collections are atomic!

6

Keys Myth

- There are Natural and Surrogate key
- WRONG!
- There are only keys
 - Who defines what is natural, and what is not natural?
 - Slightly better definition: “A key is natural if it is used outside the database and application, in real life” (definition by Fabian Pascal, corrected by Dejan Sarka)

7

Natural Keys Myth

- If you do not have a natural key, you do not have a set
- WRONG!
- Just give me a set, any set!
- A set is purely logical and abstract term
 - The one who defines the set, defines what makes elements unique as well
 - Also, do not forget on abstraction

8

Redundancy Myth

- Normalization is about eliminating redundancy
- Partially wrong
- Normalization is about eliminating redundancy and incompleteness

9

SQL Language Myth

- SQL language is essential for relational databases
- WRONG!
- The question of using relational model in a database is orthogonal to the question which language is supported by the RDBMS
 - The “D” language proposal (Date, Darwen)
 - Query By Example (QBE)

10

Truncate Table – DDL or DML?

- Wikipedia: “In SQL, the TRUNCATE TABLE statement is a Data Definition Language (DDL) operation that marks the extents of a table for deallocation (empty for reuse)”
 - So it is DDL?
- WRONG!
- However: “The TRUNCATE TABLE mytable statement is logically (though not physically) equivalent to the DELETE FROM mytable statement (without a WHERE clause)”
 - In a relational model, we are manipulating with data on the logical level; the physical implementation is left to the database management system
- So TRUNCATE TABLE is DML!

11


ALTER TABLE SWITCH – DDL or DML?

- Syntax is DDL
 - SQL Server changes system pages
 - So it is DDL?
- WRONG!
- However, logically, you just move the data from one table or partition to another table or partition
 - You do not change the schema at all
- Therefore, semantically, this is a DML statement

12

Set Operators (1)

- UNION, INTERSECT and EXCEPT operators are commonly called Set Operators
 - Represented by Venn diagrams



- WRONG!

13


Set Operators (2)

- Set operators work on sets and produce a set
- Relational operators work on relations and produce a relation; SQL operators UNION, INTERSECT and EXCEPT produce relations
- So they are relational operators!

14

Set Operators (3)

- Need to change the presentation!



15

Denormalization Myth

- Star schema is denormalized
- Not necessary
- Dimensions and fact tables
 - Dimension are denormalized only if you have rigid attribute relationships in natural hierarchies
 - Transaction fact tables are denormalized only if you include degenerate dimensions

16

Fact Table Myth

- Fact table is called fact table because it holds facts
- Yes, but... but... but...
- Don't all tables hold facts?
 - Fact tables should have different name
 - Central tables, for example

17

SCD Myth

- There are three solutions for the SCD problem (type 1, type 2, type 3)
- WRONG!
- There are more and better solutions
 - Create another dimension from the type 2 attributes, and leave them in the original dimension as well
 - Scalable solution that encompasses both type 1 and type 2

18

Class-Table Myth

- A class is the same thing as a table
- WRONG!
- A class is the same thing as a domain (data type)
 - The definitions are the same
 - Relational model is not limited to “simple” types
 - Relational model allows “typed” tables

19

Table-Class Myth

- A table is the same thing as a class
- WRONG!
- A table is a variable
 - It behaves like a variable (logically, of course)
 - Update means replacing on set with another
 - 1st Zermelo - Fraenkel axiom

20

Subtype Myth

- A subtype is the same thing as a subclass
- WRONG!
- What is inherited by a subclass?
- What is inherited by a subtype?


21

Myth above All Myths

- In this room, there are people who can drink more beer and schnapps than your speaker
- Maybeeee.....
- But it has to be proven!

22

Q & A



• Thank you!

- Only for SQL Grillen attendees: 40% discount on e-book!
 - Discount Code: **SQL2016DG40**
 - Start date: June 1st 2017
 - Expiry date: June 30th 2017
 - Create a login on the Packt site www.packtpub.com and add the book to the cart
 - Click "View Cart"
 - "Do you have a promo code?" field enter (code provided above)
 - Click the "Apply" button to apply the discount